

A Materialised Three-Way Merge for AegisSheet

Exploratory research note

FP Launchpad

2 September 2026

Abstract

The verified AegisSheet model stores the whole event history and merges by union; each event carries the timestamp set of its causal past, so state is quadratic in history length. This note asks whether the ancestor argument of a three-way merge can replace that metadata. We define a materialised state with a componentwise three-way merge and test it differentially against the existing model on random version DAGs. No counterexample appears in 3000 executions when a removal names the tokens it kills; clearing tokens instead converges but breaks commutation. Three hand-derived witnesses, checked by the executable reference, show that the last position of a removed row or column cannot be dropped while a concurrent revival is possible, so the retention residue is one register per axis identifier until removal is causally stable. The measured state is linear rather than quadratic in history. Nothing here is machine-checked; the mechanization target is a product Join proof.

1 Goal, hypotheses, and status

Goal. Re-encode AegisSheet as a genuine three-way MRDT with a materialised state so that the lowest-common-ancestor argument of `merge` replaces the per-event causal set, the append-only event log, and, where possible, the roster-acknowledged purge protocol, while preserving every observation of the current model.

Hypotheses and falsifiers.

- H1** Cells and axes forget. A state of live keep tokens, latest-position registers, active cell versions, and active range versions, merged componentwise by the rule $\text{mvr}(l, a, b) = (l \cap a \cap b) \cup (a \setminus l) \cup (b \setminus l)$ for sets and last-writer-wins for registers, yields the same observation as the union model at every reachable version. Falsifier: a version DAG on which the observations differ, two merge topologies over one event set that disagree, or a merge that differs from replaying the union.
- H2** The latest position of a removed axis identifier can be dropped once no live range names it. Falsifier: two histories with identical live sheets whose continuations require different positions for the identifier.
- H3** Purge dissolves into ordinary version deletion once state is materialised. Falsifier: a late concurrent write that revives content the paper requires to stay purged.

Status. Table 1 uses the evidence vocabulary conjectured, refuted, machine-checked, validated, measured, specified, and staged. Nothing in this note is machine-checked. The Python reference is validated against the Lean model only on the twenty SPOT fixtures listed in Section 9.

Claim	Status	Evidence
H1, named removal	validated; conjectured in general	no counterexample in 3000 executions, 117,290 versions, 21,375 registered-ancestor and 851 virtual-base merges, three seeds; Lean proof owed
H1, clearing removal	validated for observation; refuted for replay	converges and matches the reference; 103 of 3000 executions have a merge that differs from timestamp-order replay
H2	refuted	three hand-derived witnesses checked by the reference (Section 5); 16 to 55 failing executions per 1000 in the harness
H3	staged	purge is not modelled by the harness
State growth	measured	Table 3: reference quadratic, materialised linear
Undo revives a later removal	finding, reproduced	Section 6; a property of the current model

Table 1: Status of each claim.

Trusted boundary. The reference semantics is the executable Lean model in `Sal/MRDTs/Instances/AegisSheet.lean`, which the repository validates against Tables 3 and 4 of the AegisSheet paper. The Python transcription of that model is trusted only to the extent of the fixture check. The candidate design is specified here and implemented in Python; it has no Lean declaration.

2 The reference model

The displays below transcribe the Lean definitions named in each definition’s title. Purge markers are omitted throughout (labelled compression); the harness does not generate them.

Definition 2.1 (events and state; `SheetOp`, `D`). *An event is $e = (t, r, \langle \text{seen}, \text{cmd} \rangle)$ with Lamport timestamp t , replica r , a finite set `seen` of timestamps, and a command that is either a direct action or `undo( $t'$ ,  $\text{inv}$ )` carrying an inverse action. The effect `action( $e$ )` is the direct action or inv . Actions are*

$\text{axis}(\text{kind}, \alpha, \text{id}, \text{before}, \text{after})$ $\alpha \in \{\text{row}, \text{col}\}, \text{kind} \in \{\text{insert}, \text{move}, \text{remove}, \text{restore}\},$
 $\text{cell}(\rho, \gamma, \text{before}, \text{after}, \text{overwrites})$ $\text{before}, \text{after}$ finite value sets, overwrites timestamps,
 $\text{range}(\text{id}, \text{before}, \text{after}, \text{overwrites})$ $\text{before}, \text{after} \in \text{Option}(\text{RangeSpec}).$

The MRDT state is a finite set E of events; $\text{do}(E, e) = E \cup \{e\}$; $\text{merge}(_, a, b) = a \cup b$. The ancestor argument is unused.

Definition 2.2 (axis liveness; `keepsAxis`, `removesAxis`, `axisLive`). *An event e keeps (α, id) when it is an axis update of that identifier with $\text{after} \neq \perp$, or a cell update whose α -coordinate is id . It removes (α, id) when it is an axis update of that identifier with $\text{after} = \perp$.*

$\text{keepTimes}_E(\alpha, \text{id}) = \{t(e) \mid e \in E, e \text{ keeps } (\alpha, \text{id})\},$
 $\text{removed}_E(\alpha, \text{id}, \tau) \iff \exists e \in E. e \text{ removes } (\alpha, \text{id}) \wedge \tau \in \text{seen}(e),$
 $\text{live}_E(\alpha, \text{id}) \iff \text{known}_E(\alpha, \text{id}) \wedge \exists \tau \in \text{keepTimes}_E(\alpha, \text{id}). \neg \text{removed}_E(\alpha, \text{id}, \tau),$

where known_E holds when some axis update of the identifier is in E . A removal defeats exactly the keep tokens it observed; a keep it did not observe survives it.

Definition 2.3 (positions, cells, ranges; `axisPositions`, `cellValues`, `rangeValues`). A candidate for (α, id) is an axis update with $after \neq \perp$; $positions_E(\alpha, id)$ is the set of after values of candidates with no later candidate, a singleton or empty. A cell version e at (ρ, γ) is overwritten when some cell update in E lists $t(e)$ in its overwrites; $cells_E(\rho, \gamma)$ is the union of after over non-overwritten versions when both axes are live, and $\emptyset$ otherwise. Range versions are analogous, and $ranges_E(id)$ collects their non- $\perp$ after specifications.

Definition 2.4 (range resolution; `resolveRange`). For a specification $(\rho_1, \rho_2, \gamma_1, \gamma_2)$, a live endpoint resolves to itself. A dead first endpoint resolves to the live identifier at the least live position greater than the endpoint's last position; a dead last endpoint to the live identifier at the greatest live position less than it. Resolution fails when an endpoint has no position or the resolved rectangle is inverted.

Definition 2.5 (issuance; `applicableB`, `validUndo`, `inverseFor`). An event is issuable at E when $t(e) \notin times(E)$, $seen(e) = times(E)$, $t(e)$ exceeds every timestamp in E , and its before-image and overwrites list equal the current values at E (compressed: the per-kind clauses of `directApplicable`). An undo is issuable when its target is in $seen(e)$, was issued by the same replica, and the carried inverse equals `inverseFor` of the target, which swaps before and after and sets $overwrites = \{t(target)\}$ for cells and ranges. No clause requires the axes of an undone cell write to be live.

The clause $seen(e) = times(E)$ is the cost this note attacks: every event stores the timestamp set of its entire causal past.

3 The candidate design

Definition 3.1 (materialised state).

$known : \{row, col\} \rightarrow \mathcal{P}(Id)$	<i>grow-only identifier sets,</i>
$tok : (\{row, col\} \times Id) \rightarrow \mathcal{P}(Ts)$	<i>live keep tokens,</i>
$pos : (\{row, col\} \times Id) \rightarrow Ts \times Position$	<i>latest position register,</i>
$cells : (Id \times Id) \rightarrow \mathcal{P}(Ts \times \mathcal{P}(Value))$	<i>active cell versions,</i>
$ranges : Id \rightarrow \mathcal{P}(Ts \times Option(RangeSpec))$	<i>active range versions.</i>

A removal carries kills, the live tokens of its identifier at issuance (the named variant); the clearing variant carries nothing and empties the token set.

Definition 3.2 (update).

$$\begin{aligned}
axis(_, \alpha, id, _, p \neq \perp) : tok(\alpha, id) += \{t\}, \quad pos(\alpha, id) &:= \max_{ts} (pos(\alpha, id), (t, p)), \quad known += id; \\
axis(_, \alpha, id, _, \perp) : tok(\alpha, id) -= kills \quad (clearing: tok(\alpha, id) &:= \emptyset); \\
cell(\rho, \gamma, _, v, ov) : tok(row, \rho) += \{t\}, tok(col, \gamma) += \{t\}, \\
cells(\rho, \gamma) &:= \{x \in cells(\rho, \gamma) \mid ts(x) \notin ov\} \cup \{(t, v)\}; \\
range(id, _, s, ov) : ranges(id) &:= \{x \in ranges(id) \mid ts(x) \notin ov\} \\
&\cup \{(t, s)\}.
\end{aligned}$$

Definition 3.3 (merge). For $\text{mvr}(l, a, b) = (l \cap a \cap b) \cup (a \setminus l) \cup (b \setminus l)$,

$$\begin{aligned} \text{merge}(l, a, b).\text{known} &= l.\text{known} \cup a.\text{known} \cup b.\text{known}, \\ \text{merge}(l, a, b).\text{tok}(k) &= \text{mvr}(l.\text{tok}(k), a.\text{tok}(k), b.\text{tok}(k)), \\ \text{merge}(l, a, b).\text{pos}(k) &= \max_{\text{ts}} \{x.\text{pos}(k) \mid x \in \{l, a, b\}, k \in \text{dom}(x.\text{pos})\}, \\ \text{merge}(l, a, b).\text{cells}(k) &= \text{mvr}(\dots), \quad \text{merge}(l, a, b).\text{ranges}(k) = \text{mvr}(\dots). \end{aligned}$$

The observation is computed as in Section 2 with $\text{live}(\alpha, id) \Leftrightarrow id \in \text{known}(\alpha) \wedge \text{tok}(\alpha, id) \neq \emptyset$ and the register in place of the latest candidate.

The design has no cross-component rule. The reference’s update-wins policy (a concurrent write or move defeats a removal) is already token-based: the write’s token is not in the ancestor, so mvr keeps it. Every component is an observed-remove set or a last-writer-wins register.

4 H1: differential campaign

Conjecture 4.1 (H1). For every legal execution of the version-DAG semantics with the canonical virtual merge base, the materialised design with named removal has the same observation as the union model at every version.

Method. The harness (`model.py`) runs the Python reference and each materialised variant in lockstep on random version DAGs. Operations are generated from the reference state so that before-images, overwrite lists, and killed-token lists are those the issuance predicate demands; identifiers and timestamps are fresh by construction. Two to four replicas perform one to three operations each per round for three to eight rounds; a merge between two heads occurs with probability 0.7 per round; forced convergence rounds then merge all heads along different topologies. When a recorded version’s event set equals the heads’ intersection it is the ancestor; otherwise the ancestor is the fold of the intersection in timestamp order, which is the canonical virtual merge base. At every version the harness compares observations, compares observations of versions reached by different topologies over the same event set, and compares the merge result with the timestamp-order replay of the union.

Result. Table 2 reports three seeds of 1000 executions. The named-removal design has no failing execution. The clearing variant matches the reference and converges but fails the replay comparison: a removal that clears all tokens and a concurrent keep do not commute, so replay order matters and a Lean proof would need the arbitration arm of the replay policy. The named variant makes every pair of operations commute, which is the route every existing production instance uses.

5 H2: the last position of a dead axis is load-bearing

Each witness is hand-derived from the definitions of Section 2 and checked by evaluating the Python reference; the checks are the functions `h2_witnesses` and `h2_pair3`. Identifiers: rows r_0, r_1, r_2 at positions 10, 20, 30, column c_0 .

Counterexample 5.1 (range resolution). Let L insert r_0, r_1, r_2, c_0 and a range (r_0, r_2, c_0, c_0) . History W_1 removes r_0 . History W_2 moves r_0 to position 25 and then removes it. Both leave

Design	seed 1	seed 2	seed 3	first failure class
keep registers, named removal	0	0	0	
keep registers, clearing removal	27	31	45	merge $\neq$ replay
drop unreferenced dead registers, named	28	23	16	position of revived axis
drop all dead registers, named	10	10	7	range resolution
ancestor ignored (control)	186	173	146	cells, ranges
eager range re-anchoring (control)	634	642	683	ranges

Table 2: Failing executions out of 1000 per seed. Each seed produced about 39,000 versions, 31,500 operations, 7,100 merges at a registered ancestor, and 280 merges at a virtual base.

live rows $\{r_1, r_2\}$ at unchanged positions and the same range specification. The reference resolves the range to (r_1, r_2, c_0, c_0) after W_1 and to (r_2, r_2, c_0, c_0) after W_2 . A state that forgets r_0 's last position cannot give both answers.

Counterexample 5.2 (eager re-anchoring). *From the same L , branch A removes r_0 and branch B concurrently inserts r_3 at position 15. The reference resolves the merged range to (r_3, r_2, c_0, c_0) . Rewriting the range to r_0 's successor at removal time gives (r_1, r_2, c_0, c_0) . Lazy resolution over the dead position is part of the specification, not an implementation choice.*

Counterexample 5.3 (update-wins revival). *Let L insert r_0 at 10 and c_0 , and write cell (r_0, c_0) . Branch A_1 moves r_0 to 52 and then removes it; branch A_2 only removes it. Both branch states have identical live sheets: r_0 is dead. Branch B concurrently writes (r_0, c_0) . In both merges the write's token revives r_0 ; the reference places it at 52 after A_1 and at 10 after A_2 . The merge receives identical L and S_B , so if $S_{A_1} = S_{A_2}$ it cannot answer both. The information lives only on the removing branch.*

The third counterexample does not involve ranges or undo, and it survives restricting undo to live axes (55 failing executions of 1000 at seed 7). The retention residue is therefore: one position register per axis identifier that has ever been known, retained until the removal is causally stable, that is, until no replica can still issue an operation concurrent with it. Under update-wins this is unavoidable; under a remove-wins policy the register could be dropped at removal.

6 Finding: undo revives a causally later removal

Finding 6.1. *In the current model, with row r_0 and column c_0 and a write at $t = 3$, removing r_0 at $t = 4$ and then undoing the $t = 3$ write at $t = 5$ is a legal history on one replica: the undo's target is observed and was issued by the same replica, and no issuance clause requires the axes to be live. The undo is a cell event and therefore a keep token for r_0 with a timestamp the removal did not observe. The row becomes live again at position 10 with an empty cell. Reproduced by `undo_revival_witness`.*

Table 4 of the paper covers undo against a *concurrent* removal; it does not say what undoing an edit should do to a row removed causally later. The behaviour is a consequence of representing undo as a cell event. If it is unintended, the fix is an issuance clause (undo of a cell write requires live axes) or a non-keeping inverse; either changes the public sequential specification and should be decided before the Lean port.

7 H3: purge

The harness does not generate purge markers, so H3 is staged. Two observations shape the next experiment. First, in the materialised design a purge is an ordinary deletion of cell versions at dead coordinates, and the `mvr` rule already keeps a concurrent write’s version because it is absent from the ancestor, which is the paper’s stated intent that a late change restores the axis but not the purged contents. Second, the roster acknowledgement of the current protocol is the causal-stability condition that Section 5 identifies as the moment the dead register can be dropped. The conjecture for the next step is that the acknowledgement protocol is needed for the $O(1)$ register, not for the cell payloads.

8 Measured: state size

Rounds	Operations	Reference (timestamps stored)	Materialised, named, keep
4	23	131	21
8	46	594	42
16	94	3,048	90
32	184	13,866	169

Table 3: Mean over 20 executions with 3 replicas, seed 1, at the final version of replica 0. The reference count is one per event plus its `seen` set and overwrite list; the materialised count is tokens, registers, and versions.

Doubling the history quadruples the reference and doubles the materialised state. The gap at 184 operations is a factor of 82. These are counts of stored timestamps in the Python models, not bytes of any runtime.

9 Controls

Positive control. The Python reference reproduces twenty hand-derived expectations copied from the Lean SPOTs: update defeats concurrent removal, later removal wins, concurrent writes keep both values, selective undo keeps the remote value, move defeats removal, later move wins, insert survives adjacent removal, undo of remove restores the position, undo of insert removes only the inserted identifier, the two Table 4 move-undo cases, and seven Figure 1 range cases (`check_fixtures`).

Negative controls. Two deliberately wrong designs fail on every seed (Table 2): ignoring the ancestor, which reintroduces removed versions, and eager range re-anchoring. A harness that could not distinguish these would carry no evidence.

10 Mechanization correspondence

Display	Lean declaration	Status
Reference events, state, liveness, positions, cells, ranges, resolution, issuance	<code>Instances.AegisSheet: D, axisLive, axisPositions, cellValues, rangeValues, resolveRange, applicableB</code>	machine-checked definitions; transcription validated on 20 fixtures
Materialised state, update, merge	none	specified; Python implementation
Conjecture H1	none	validated on the campaign scope
Counterexamples 5.1 to 5.3, Finding 6.1	none	hand-derived, checked by the Python reference; to be pinned as Lean SPOTs

Lean plan. The materialised design is a product of observed-remove sets and last-writer-wins registers. With named removal every operation pair commutes, so the `Join` proof can follow the universal-equation route, or the product theorem `joinAt_prod` with a per-component `CanonicalJoinLaws` in the style of the multi-value register. The sequential certificate needs a new representation relation; `no_view_only_step` fixes that it must retain tokens, cell version identities, and range version identities, which the materialised state does.

11 Open obligations

1. Decide the undo-revival question of Section 6 before porting; it changes the sequential specification.
2. Model purge in the harness and test H3: purge as version deletion, concurrent writes surviving, and the register-dropping rule at causal stability.
3. Pin Counterexamples 5.1 to 5.3 and Finding 6.1 as Lean SPOTs against the existing model, with PASS and FAIL companions.
4. Port the named-removal design to Lean as a second package beside the current one; prove `Join` and the sequential certificate; keep the union model until the runtime measurement is in.
5. Extend the harness to merges between non-head versions and to undo of undo; report the generator's coverage.