

A Verification Framework for State-Dependent MRDTs: describe only the sensible operations

Sal metatheory note — framework draft
FP Launchpad, IIT Madras

July 4, 2026

Abstract

Mergeable replicated data types (MRDTs) are verified by a standard division of labour: the data-type author discharges a fixed catalogue of local, equational *verification conditions* (VCs), and a once-and-for-all soundness meta-theorem lifts them to RA-linearizability of every reachable configuration. Every published catalogue quantifies its VCs *over all states*. This is adequate for *flat* data types — sets, counters, registers — whose operations are total: meaningful on every state. It is *inadequate* for *state-dependent* data types, whose operations are only meaningful on well-formed states and whose effect depends on the state they observe. The sharp cases are tombstone-free sequence and rich-text CRDTs (a physical-delete RGA; Peritext without tombstones): deletion physically removes a node and *re-anchors* its dependents, so an operation’s recorded position can become stale, and its commutation with a concurrent operation holds only on reachable states. Whole-state VCs are simply *false* on the infeasible states such a type admits.

We present a *conditioned* verification framework. The author enriches the signature with two predicates — an *applicability* predicate `app` (“when is this operation sensible?”) and a reachability invariant `lnv` — and discharges VCs *conditioned* on them; the framework proves RA-linearizability once and for all. The developer thus *describes only the sensible operations* (as the RGA already does with insertion anchors) and is never asked to reason about infeasible states. The technical core is an *applicability-aware* linearization order together with a commutation discipline — path-faithfulness plus causal freshness — under which state-dependent, re-anchoring operations provably commute. The development is mechanized in Lean 4; we state which results are kernel-checked and which are in progress.

1 Introduction

An MRDT is a replicated data type whose replicas fork, apply operations locally, and reconcile divergence by a three-way *merge* `merge(l, a, b)` relative to the lowest common ancestor l of the two divergent states a, b . Correctness is *RA-linearizability*: every version in the store is explained by a sequencing of exactly the events it has observed, consistent with a linearization order lo derived from visibility and a conflict-resolution relation rc .

The prevailing methodology reduces this global property to a fixed catalogue of *local, equational* VCs over `do`, `merge` and `rc`, discharged per data type, and a soundness meta-theorem, proved once, that lifts the VCs to RA-linearizability of every reachable configuration. Crucially, the VCs are stated as universally quantified equations *over all states* $\sigma \in \Sigma$.

The whole-state assumption, and where it breaks. Universal quantification is benign when every operation is *total* — meaningful on every state — and *state-independent* in the sense that a VC true on the reachable states is true on all states. Flat data types are of this kind: a grow-only set’s `add` inserts on any state; an OR-set’s `rem` deletes whatever tags it sees. But consider a *tombstone-free* sequence CRDT (an RGA that physically deletes):

- An insertion records the *position* it attaches to (a path to an anchor node). If a concurrent deletion has removed that anchor, the insertion must *climb* to the nearest surviving ancestor — its effect *depends on the state*.
- A deletion physically removes a node and *re-anchors* (rehomes) its children to the deleted node’s parent.
- Consequently two operations commute only on *well-formed* states: replayed on a syntactically-constructible but unreachable state, the equation is false.

A whole-state VC catalogue cannot express such a type: its commutation VC is false as a $\forall\sigma$ statement, even though it holds on every state the system can actually reach. This is not a peculiarity of the RGA; it is the generic situation for any data type with *structural* state and *positional* operations. Peritext without tombstones is the same phenomenon one level up: a formatting mark anchors to character positions, and deleting a character re-anchors the mark.

Thesis. Whole-state VC catalogues *over-specify*: they demand equations on states no execution reaches. The remedy is not to invent new equations but to *condition* the existing ones on what is reachable and sensible, and to re-prove the meta-theorem to consume the conditioned forms. The author’s added burden should be exactly the *description of sensible operations* — which, for the RGA, is already implicit in the anchor discipline — and nothing more.

Contributions.

1. A *conditioned signature* (§3) carrying an applicability predicate `app` and a reachability invariant `Inv`, with a *collapse* property: flat types take `app = Inv =  $\top$` and recover the ordinary framework verbatim.
2. A *conditioned VC discipline* (§5): commutation and the linearization order are relativized to `app` and `Inv`, so the author discharges VCs *assuming sensibility*.
3. The technical core (§6): an *applicability-aware* linearization order and a commutation discipline — *path-faithfulness* and *causal freshness* — under which state-dependent, re-anchoring operations commute; and the convergence theorem this yields.
4. A once-and-for-all *adequacy* theorem (§7) reducing RA-linearizability to the conditioned VCs, with an explicit division between the author’s obligations and the framework’s generic proof.
5. Instantiation (§8): the tombstone-free RGA as the driving instance, and Peritext-without-tombstones as an instance of the same *anchor-carrying* class that reuses the commutation discipline.

2 The replicated model and RA-linearizability

We recall the model, briefly, to fix notation; the framework of this note is orthogonal to its details.

An *event* is a triple $e = (t, r, o)$: a globally unique timestamp t , an origin replica r , and an operation o . A *configuration* carries a *store*: a directed acyclic graph of *versions*, each holding a state $\sigma \in \Sigma$ and the set $L(v)$ of events that produced it, with per-replica heads. A transition forks a replica, applies a fresh event ($\sigma \mapsto \text{do}(\sigma, e)$; we write $e(\sigma)$), queries, or **merges**: it picks two heads v_1, v_2 , finds their lowest common ancestor v_ℓ , and installs a new version with state $\text{merge}(\sigma_{v_\ell}, \sigma_{v_1}, \sigma_{v_2})$ and event set $L(v_1) \cup L(v_2)$. Visibility $\xrightarrow{\text{vis}}$ records what each event's issuing replica had seen; $e_1 \parallel e_2$ denotes concurrency ($\neg(e_1 \xrightarrow{\text{vis}} e_2)$ and $\neg(e_2 \xrightarrow{\text{vis}} e_1)$). Two events *commute*, $e_1 \rightleftharpoons e_2$, when $e_1(e_2(\sigma)) = e_2(e_1(\sigma))$ for all σ .

Definition 2.1 (RA-linearizability, per version). *A list $\pi = e^1 \dots e^n$ witnesses a version holding state s with event set E if (i) π enumerates E without repetition; (ii) π extends the linearization order lo ; and (iii) it folds to the state, $e^n(\dots e^1(\sigma_0) \dots) = s$. A configuration is RA-linearizable if every version in its store admits a witness.*

The order lo places e_1 before e_2 when either $e_1 \xrightarrow{\text{vis}} e_2$ and $e_1 \not\rightleftharpoons e_2$, or $e_1 \parallel e_2$, $e_1 \not\rightleftharpoons e_2$, rc orders them, and e_2 is not already overwritten by a later non-commuting event. We work throughout with the *set-relative* refinement lo^E of this order, in which the overwriter clause ranges over the event set E of the version being linearized rather than the whole configuration; this refinement is forced (an event can gain overwriters from a concurrent branch) and $\text{lo} \subseteq \text{lo}^E$ pointwise, so every lo^E -witness is an lo -witness.

3 The conditioned signature

Definition 3.1 (conditioned MRDT signature). *A conditioned MRDT signature is a tuple*

$$\mathcal{D} = \langle \Sigma, \sigma_0, \text{do}, \text{merge}, \text{rc}, \approx, \text{Inv}, \text{app} \rangle$$

extending an ordinary MRDT signature $\langle \Sigma, \sigma_0, \text{do}, \text{merge}, \text{rc} \rangle$ with

- *an observational equivalence $\approx \subseteq \Sigma \times \Sigma$ — an equivalence relation that is a congruence for do and merge (Def. 3.2); on flat types it may be taken as structural equality;*
- *a reachability invariant $\text{Inv} : \Sigma \rightarrow \text{Prop}$ — an over-approximation of the reachable states;*
- *an applicability predicate $\text{app} : \text{Op} \times \Sigma \rightarrow \text{Prop}$ — “operation o is sensible on state σ .”*

A signature is flat when $\text{Inv} = \top$, $\text{app} = \top$, and $\approx$ is structural equality; then all conditioned notions below reduce definitionally to the ordinary ones.

Definition 3.2 (observational equivalence). *$\sigma_1 \approx \sigma_2$ iff every query supported by the data type returns the same value on σ_1 and σ_2 . Convergence is stated up to $\approx$: two replicas that have observed the same events must be observationally equivalent, not necessarily structurally equal.*¹

¹This is the notion of *convergence modulo observable behaviour*: two states may differ in representation (a balanced vs. unbalanced search tree; a physically-deleted node left as unreachable junk) yet agree on every observation. For the tombstone-free RGA it turns out that reachable convergence holds even *structurally*, so $\approx$ may be taken as equality; we keep it abstract because state-dependent types in general need it.

Running example: the tombstone-free RGA. The state is a forest over node identifiers: Σ maps each live identifier t to a payload and an anchor $\text{anc}(t)$ (its parent), with a root sentinel 0. We write $\text{live}(\sigma, t)$ for “ t is present in σ ”. Two operations:

$$\begin{aligned} \text{Ins}(e, p, a) : & \quad \text{do}(\sigma, \text{Ins}(e, p, a)) = \sigma[t \mapsto (e, \text{resolve}(\sigma, a:p))], \\ \text{Del}(p, x) : & \quad \text{rehome each child of } x \text{ to } \text{resolve}(\sigma, p), \text{ then remove } x, \end{aligned}$$

where $\text{resolve}(\sigma, L)$ returns the first identifier of the list L that is live in σ , or the root 0 if none is (the *climb*). An insertion records a *path* p and an *anchor* a ; a deletion records the path p to which its children rehome. The applicability and invariant predicates capture “sensible”:

$$\begin{aligned} \text{app}(\text{Ins}(e, p, a), \sigma) & \equiv (a = 0 \wedge p = []) \vee (\text{live}(\sigma, a) \wedge p \text{ is the true ancestor chain of } a), \\ \text{app}(\text{Del}(p, x), \sigma) & \equiv \text{analogously, } x \text{ live and } p \text{ its true rehome path,} \\ \text{Inv}(\sigma) & \equiv \text{wf}(\sigma) \wedge \text{idmono}(\sigma), \end{aligned}$$

where wf says the forest is well-formed (every anchor live or the root) and idmono says anchors strictly decrease ($\text{anc}(\sigma, t) = 0 \vee \text{anc}(\sigma, t) < t$), a consequence of monotone timestamp allocation. The anchor discipline the RGA already carries *is* its applicability predicate; the author writes nothing new to describe sensible operations.

Example 3.3 (the unconditioned commutation VC is false). *This is why conditioning is necessary, not merely convenient. Take two insertions issued concurrently by different replicas: a records anchor 3 and path [7] (it creates a node under 3, climbing the path if 3 is gone), and b creates a node with identifier 7 at the root. Evaluate the two orders on the initial root-only state σ_0 :*

- $\text{do}(\text{do}(\sigma_0, a), b)$ — *a first: it resolves its list [3, 7] against σ_0 , where both 3 and 7 are absent, and climbs to the root, anchoring its node at 0; then b creates node 7.*
- $\text{do}(\text{do}(\sigma_0, b), a)$ — *b first: it creates node 7; now a resolves [3, 7] with 3 absent but 7 live, anchoring its node at 7.*

The two results disagree on the parent of a’s node (0 versus 7): the whole-state commutation equation

$$\forall \sigma. \quad \text{do}(\text{do}(\sigma, a), b) = \text{do}(\text{do}(\sigma, b), a)$$

*is therefore false, already at $\sigma = \sigma_0$ (**naive_general_swap_false**). No execution reaches this counterexample: a records an anchor 3 that is not present and a path naming b’s identifier 7, yet a and b are concurrent — a’s replica never observed b, so it could neither anchor to an absent node nor name b’s fresh identifier. These are exactly the tuples **app** (and causal freshness, §6) excludes; restricted to applicable tuples, the two insertions do commute (Lemma 6.5). A whole-state catalogue quantifies its commutation VC over $\sigma = \sigma_0$ with this a and so cannot state a true VC for the type at all; the conditioned catalogue can. The RGA already encodes the fix — its shipped commutation VC is gated on $\text{accurate} \wedge \text{fresh} \wedge \text{wf}$ precisely because the ungated equation is this counterexample.*

4 Sensible operations and reachability

The two predicates play complementary roles, and each carries exactly one obligation.

Definition 4.1 (reachability invariant). *Inv is a reachability invariant if $\text{Inv}(\sigma_0)$ holds and Inv is preserved by every transition: for all σ with $\text{Inv}(\sigma)$ and every applicable event e , $\text{Inv}(e(\sigma))$; and for all Inv-states l, a, b arising as an honest LCA-and-branches triple, $\text{Inv}(\text{merge}(l, a, b))$. We write Reach for the states satisfying Inv.*

Establishing Inv as a reachability invariant is the author’s obligation; it is a routine induction over do and merge (for the RGA, preservation of wf under merge requires exactly idmono, which the same induction supplies). No other whole-state reasoning is asked of the author.

Remark 4.2 (the developer-effort ledger). *The essential/author column and the generic/framework column are:*

<i>Author provides (essential)</i>	<i>Framework proves (once)</i>
app — spec of sensible operations	convergence of concurrent operations
Inv + reachability-invariance (Def. 4.1)	existence of canonical states $\sigma(E)$
conditioned VCs (§5, assuming app)	the linearization machinery
	RA-linearizability (Thm. 7.2)

The author’s column is the specification — one cannot verify a state-dependent type without saying which states are valid and which operations sensible — and no more. The failure mode of a whole-state framework is that it leaks framework-column work into the author’s lap; §6 shows how the conditioned framework keeps the columns separate.

5 The conditioned verification conditions

The single primitive is conditioned commutation.

Definition 5.1 (conditioned commutation). *Operations o_1, o_2 commute where sensible, written $o_1 \rightleftharpoons_{\text{app}} o_2$, if for every σ with $\text{Inv}(\sigma)$, $\text{app}(o_1, \sigma)$ and $\text{app}(o_2, \sigma)$,*

$$\text{do}(\text{do}(\sigma, o_1), o_2) \approx \text{do}(\text{do}(\sigma, o_2), o_1).$$

Commutation is demanded only at reachable states where both operations are sensible, and only up to observation. The relational VCs on rc are the usual ones with $\rightleftharpoons$ replaced by $\rightleftharpoons_{\text{app}}$: RC-NON-COMM (non-commuting distinct-replica operations are rc-ordered in exactly one direction), NO-RC-CHAIN (no $o_1 \xrightarrow{\text{rc}} o_2 \xrightarrow{\text{rc}} o_3$), and a conditioned COND-COMM (a resolved conflict is invisible once an overwriter runs). These are the author’s update-layer obligations, discharged *assuming* applicability of the operations involved — which, for state-dependent types, is what makes them provable at all.

6 Canonical states and convergence

This section is the technical core: how conditioned VCs yield a well-defined canonical state per event set, and the commutation discipline that makes it work for re-anchoring operations.

6.1 The applicability-aware order

The set-relative order lo^E (§2) orders *non-commuting* pairs. For state-dependent operations a second source of ordering appears: one operation’s *applicability* may depend on another. A concurrent deletion can render an insertion’s recorded anchor stale; the insertion still commutes with the deletion (it climbs), but the two are related by a *generation dependency* that the order must record, lest a reordering apply an operation at a state where it is neither sensible nor a harmless no-op.

Definition 6.1 (generation dependency; applicability-aware order). *Event e_2 depends on e_1 , written $e_1 \prec_{\text{app}} e_2$, if e_1 can change whether e_2 is applicable: $\exists \sigma. \text{app}(e_2, \sigma) \neq \text{app}(e_2, e_1(\sigma))$. The applicability-aware order is*

$$\text{lo}_{\text{app}}^E = \text{lo}^E \cup \{(e_1, e_2) : e_1 \xrightarrow{\text{vis}} e_2 \wedge e_1 \prec_{\text{app}} e_2\}.$$

For the RGA, $\prec_{\text{app}}$ is the syntactic fact “ e_2 ’s recorded path names e_1 ’s identifier”; it is discharged directly from the operation shapes.

Definition 6.2 (no-op-feasible enumeration). *A list $\pi = e^1 \cdots e^n$ is no-op-feasible if at each prefix fold $\sigma_i = e^i(\cdots e^1(\sigma_0) \cdots)$ the next event is either applicable at σ_i or acts as the identity there ($e^{i+1}(\sigma_i) \approx \sigma_i$).*

Two feasibility repairs are needed and are *complementary*, not alternatives: lo_{app}^E excludes reorderings that violate a generation dependency (which mere no-op tolerance would wrongly admit), while no-op tolerance admits the harmless *redundant* reorderings that strict applicability would wrongly reject (e.g. two concurrent deletions of one node: whichever runs second is a no-op, so no fully-applicable enumeration exists, yet both orders converge). The order handles the first; no-op-feasibility the second.

6.2 The commutation discipline for anchor-carrying types

The convergence proof reorders an enumeration to a canonical form by adjacent swaps of lo_{app}^E -incomparable events. Each swap must hold *at the prefix fold where the pair sits* — a state at which either event may be stale, because prior events in the prefix may have deleted its anchor. The discipline below identifies exactly what makes such a swap valid without demanding either event be applicable there.

Definition 6.3 (path-faithfulness). *An operation is faithful at σ if its recorded list tracks the true tree of σ above its climb target: for an insertion with list L , writing $c = \text{resolve}(\sigma, L)$, if c is live then $\text{resolve}(\sigma, L \setminus \{c\}) = \text{anc}(\sigma, c)$ (the list continues up the real ancestor chain); a deletion is faithful if its rehome path likewise resolves to its target’s true parent. Write $\text{Faithful}(o, \sigma)$.*

Definition 6.4 (causal freshness). $\text{NoFreshClash}(a, b)$ holds if b ’s (fresh) identifier does not occur in a ’s recorded list — the condition a real execution guarantees between concurrent events, since neither observed the other and identifiers are allocated afresh.

Faithfulness is strictly weaker than applicability ($\text{app}(o, \sigma) \Rightarrow \text{Faithful}(o, \sigma)$) and, unlike applicability, is *stable under re-anchoring*: a deletion that rehomes nodes preserves the faithfulness of a concurrent operation’s recorded path, because rehoming updates the true tree consistently with the climb. This is the property that lets the swap fire at stale states.

Lemma 6.5 (faithful commutation). *Let σ satisfy `Inv` with root absent, and let a, b be distinct-timestamp events with `Faithful`(a, σ), `Faithful`(b, σ), `NoFreshClash`(a, b) and `NoFreshClash`(b, a). Then*

$$\text{do}(\text{do}(\sigma, a), b) \approx \text{do}(\text{do}(\sigma, b), a).$$

general_swap_bothFaithful

Lemma 6.5 is mechanized (kernel-checked, all four insertion/deletion combinations). It is the reusable heart of the framework for anchor-carrying types: it says two re-anchoring operations commute whenever each is merely faithful and they do not name each other’s fresh identifiers — neither need be applicable. The asymmetric special case (b applicable, a faithful) is `general_swap`; that the symmetric case also holds is what guarantees the convergence swap fires *at every* prefix fold, where both events may be stale.

6.3 Convergence

Faithfulness must hold not merely at one state but at every prefix fold the canonicalization visits. This is a reachability-style invariant — call it *faithful-threading* — and it is where the framework absorbs the work that a whole-state approach would push onto the author.

Proposition 6.6 (faithful-threading, target form). *Along any lo_{app}^E -respecting, no-op-feasible enumeration of a backward-closed E , every event is faithful at the prefix fold at which it is applied.*

Threading is established from the author’s `app/Inv` data by a single-step preservation argument: faithfulness is preserved by an applicable insertion, and by an applicable deletion (the deletion splices its target out of the true chain, re-linking children to the target’s successor). The insertion step and the correct invariant strengthening are mechanized; the deletion step is the one remaining lemma (§9).

Theorem 6.7 (conditioned convergence, target form). *Assume the conditioned VCs (§5), Lemma 6.5, and Proposition 6.6. Then any two enumerations of a backward-closed event set E that respect lo_{app}^E and are no-op-feasible fold from σ_0 to $\approx$ -equivalent states.*

Convergence licenses the *canonical state* $\sigma(E)$ — the $\approx$ -class obtained by folding any such enumeration — and reduces “a version linearizes” to “its stored state $\approx \sigma$ of its event set,” with witness lists eliminated from all downstream reasoning. The proof is the ordinary bottom-up canonicalization (proved once for flat types), with each adjacent swap discharged by Lemma 6.5 under Proposition 6.6; nothing in it mentions `resolve`, anchors, or the RGA. That is the sense in which convergence is generic: it consumes the faithful-commutation VC (Lemma 6.5) and the threading obligation (Proposition 6.6) as hypotheses, and each data type supplies those from its own operation algebra.

7 Adequacy

Convergence is the update-layer half. The merge-layer half is a single equation between canonical states.

Definition 7.1 (Join Lemma). *For backward-closed event sets E_1, E_2 of a reachable configuration, $\sigma(E_1 \cup E_2) \approx \text{merge}(\sigma(E_1 \cap E_2), \sigma(E_1), \sigma(E_2))$.*

By the LCA property $L(v_\ell) = L(v_1) \cap L(v_2)$, the first argument is exactly the state the store hands the merge. The two halves compose once and for all:

Theorem 7.2 (adequacy). *If $\mathcal{D}$ satisfies the conditioned VCs, has Inv as a reachability invariant, and satisfies the Join Lemma, then every configuration reachable in the transition system is RA-linearizable, per version.*

The adequacy bridge is closure-indexed and proved generically: because every store version is fully causally closed, the Join Lemma applies at it regardless of the closure strength a data type declares, and the two established verification routes (a delta contract for set- and counter-shaped merges; a direct join for counter-comparison merges) are its two instances. The bridge and its instances are mechanized (`ra_linearizable3_of_joinC`), with the nine flat production MRDTs of the repository routed through it.

8 The class map and instances

Data types occupy a three-axis space: the *delta class* of the merge (group, lattice, or feasibility-conditioned), the *closure strength* the Join Lemma needs, and the *state invariant* Inv . Flat types sit on the $\text{Inv} = \top$ plane and are handled by the ordinary framework (the collapse of Def. 3.1). The conditioned framework of this note is what populates the $\text{Inv} \neq \top$ plane.

Tombstone-free RGA. The driving instance: $\text{app} =$ the anchor discipline, $\text{Inv} = \text{wf} \wedge \text{idmono}$, faithful commutation (Lemma 6.5) discharged from `resolve/rehome` algebra.

Peritext without tombstones. The same class, one level up. A formatting mark anchors to character positions as an RGA node anchors to its parent; deleting a character re-anchors the marks that span it, exactly the rehoming phenomenon. The author supplies app (a mark’s endpoints anchor live characters), Inv (well-formed marks over a well-formed sequence), and discharges the conditioned VCs by the *same* faithfulness/freshness discipline (Lemmas 6.5 and the threading of Prop. 6.6), with marks reusing the anchor machinery. It is thereby an *instance*, not a second bespoke development — which is the payoff of a framework over a per-type proof.

9 Mechanization status

The development is in Lean 4 (Mathlib; axioms `{propext, Classical.choice, Quot.sound}`; no `sorry` in the cited results). Honestly delineated:

- **Kernel-checked.** The closure-indexed adequacy bridge (`ra_linearizable3_of_joinC`) and the nine flat MRDTs routed through it; the set-relative convergence for flat types; the conditioned-convergence order reduction; and the commutation discipline for the RGA — Lemma 6.5 (`general_swap_bothFaithful`) and its asymmetric form (`general_swap`), all insertion/deletion combinations.
- **In progress.** Proposition 6.6 (faithful-threading): the correct invariant (a chain-faithfulness strengthening) and its insertion step are mechanized; the deletion-preservation step is the single remaining lemma, being discharged concurrently by direct proof and by property-based testing on reachable states. Given it, Theorem 6.7 is the ordinary canonicalization instantiated with Lemma 6.5.

- **Design/target.** Theorem 6.7 and the end-to-end hosting of the RGA (and Peritext) as conditioned instances are the framework’s intended guarantee; their assembly rests on the in-progress step above. We have verified, on reachable states, that the RGA converges even structurally, so $\approx$ may there be taken as equality — a simplification, not a soundness requirement.

Relation to prior methodologies. Whole-state VC catalogues (Neem) are the $\text{Inv} = \text{app} = \top$ collapse of this framework. The use of observational equivalence rather than structural equality follows earlier work on *convergence modulo observable behaviour* (Peepul); the contribution here is to combine it with applicability conditioning and to isolate the commutation discipline (faithfulness + freshness) that makes state-dependent, re-anchoring operations tractable, so that the author need only *describe the sensible operations*.