

Why the Path Matters

A visual account of the tombstone-free, path-carrying RGA

Sal / FP Launchpad

Abstract

The replicated growable array (RGA) in `RGA_Tombstone_Free` stores only *live* elements: a delete physically removes its target, leaving no tombstone. To stay convergent without tombstones, each operation carries the *ancestor path* of the node it references. This note explains, with pictures, two things: (i) **why the implementation converges** — the three-way merge recovers a concurrently-deleted node’s parent from the lowest common ancestor (LCA); and (ii) **why the path is necessary for the single-replica commutation proof** — to discharge the framework’s do-commutation VC, where there is no LCA, an insert whose anchor was deleted has nowhere to recover its position *unless it carries that position itself*. The qualification matters: the merge alone already converges the concurrent case *without* any path (`merge_converges_concurrent`); it is this commutation *obligation*, not operational convergence, that the path serves. We give the counterexample that diverges without the path, the fix that the path provides, and the asymmetry that makes the path mandatory for `Ins` but optional for `Del`. Every claim here is mechanised in Lean (§5).

1 The data type: a forest of live records

The state is a finite map $\sigma : \text{id} \mapsto (\text{el}, \text{anc})$ from an identity to its element and its *immediate anchor* — the identity it was inserted after. Identity 0 is the root sentinel and is never stored. So a state *is* a forest rooted at 0: an edge from a node to its anchor (Figure 1). We write $\text{ids}(\sigma)$ for the stored identities, $\text{el}(\sigma, t)$ and $\text{anc}(\sigma, t)$ for the element and anchor of t , and $\text{contains}(\sigma, t)$ for “ t is live”. Records use the reference notation (id, element, anchor).




Figure 1: The state as a forest. $\sigma = \{(1, A, 0), (2, B, 1)\}$.

Operations. There are two, and (crucially) each carries a *path* p — the ancestor chain of its referenced node, nearest first:

- `Ins(e, p, a)` at a fresh id t : insert e anchored at a . It stores $(t, e, \text{resolve}(\sigma, a :: p))$, where `resolve` returns the first *live* candidate in its list, else 0. When a is live this is just a ; when a is dead it climbs p to the nearest live ancestor.

- $\text{Del}(p, x)$: *splice* out x — physically remove it and re-parent its children to $\text{resolve}(\sigma, p)$, the nearest live ancestor of x . No tombstone is kept.

Figure 2 shows the defining move of a tombstone-free delete: the splice. Deleting 2 removes it and lifts its child 3 one level, onto 2’s parent 1.




Figure 2: The splice. A tombstone-free $\text{Del}(2)$ physically removes 2 and rehomes its child 3 to 2’s live parent 1 (red dashed = the new anchor). Compare a tombstone design, which would instead keep 2 as a dead marker.

2 Why the implementation converges: the LCA does the remembering

Concurrency is reconciled by a three-way merge $\text{merge}(L, A, B)$ where L is the least common ancestor of branches A and B . It has two steps: *survival* (an OR-set merge on the three id sets) and *re-parenting* (climb each survivor’s birth-anchor up L ’s chain). The survivor set is

$$I = (\text{ids}(L) \cap \text{ids}(A) \cap \text{ids}(B)) \cup (\text{ids}(A) \setminus \text{ids}(L)) \cup (\text{ids}(B) \setminus \text{ids}(L))$$

— an id lives iff it is kept on both sides, or freshly added on either; equivalently, a node lives unless it was in L and deleted on a side. Each survivor t takes its element $\text{el}(t)$ and *birth-anchor* $\beta(t)$ from L if present, else A , else B ; the birth-anchor is then climbed up L ’s chain — via L ’s anchor map anc_L — to the nearest live survivor:

$$\text{climb}_I(a) = \begin{cases} a, & a = 0 \text{ or } a \in I, \\ \text{climb}_I(\text{anc}_L(a)), & a \in \text{ids}(L) \setminus I. \end{cases}$$

$$\text{merge}(L, A, B) = \{ (t, \text{el}(t), \text{climb}_I(\beta(t))) \mid t \in I \}.$$

The point for convergence is the phrase “**up L ’s chain**”: even when a node was deleted on one branch, *the LCA still holds it live, with its parent*. So the merge can always recover where a survivor should attach — it reads parents from L , and never needs the operations’ paths. Figure 3 is the canonical case: A inserts 3 after 2 while B concurrently deletes 2; the merge climbs 3’s dead birth-anchor 2 up the LCA chain to the live 1.

This is the whole reason a tombstone-free RGA can converge at all: the merge has a third input, L , that retains the deleted node’s ancestry. Keep that fact in mind — the next section is about the one place where there is *no* such third input.




Figure 3: Why merge converges. 3 was born under 2 (dotted, 2’s ghost), but 2 is deleted on B . The merge climbs 3’s birth-anchor up the LCA chain — $2 \in \text{ids}(L) \setminus I$, so $\text{climb}_I(2) = \text{climb}_I(\text{anc}_L(2)) = \text{climb}_I(1) = 1$ — and rehomes 3 to 1 (red dashed). The LCA is where “2’s parent is 1” is read. **No path on the operations is used.**

3 Why the path is necessary: a single replica has no LCA

Convergence of a replicated data type also requires that operations applied in different orders on *one* replica agree — the framework discharges this as a *commutation* condition on do (single-replica application). And here is the catch: **do has no LCA**. There is only the current state. So when an insert’s anchor has already been deleted, the parent the merge would have read from L is simply *gone* — unless the insert brought it along.

The counterexample (no path). Take $\sigma = \{(1, A, 0), (2, B, 1)\}$ and the two operations $\text{Ins}(C, 2)$ (insert C as id 3 after 2) and $\text{Del}(2)$. Apply them in both orders, with a *prefix-free* insert (anchor only, no path):




Figure 4: **Divergence without the path.** The same insert/delete pair, two orders, two different trees: 3 is anchored at 1 on the left but at 0 on the right. The information “2’s parent is 1” was destroyed by the tombstone-free delete, and nothing on the right carries it. Mechanised as `prefixfree_diverges`.

The two orders disagree on 3’s anchor (1 versus 0): the design does *not* commute. Because the delete physically removed 2, the fact “ $\text{anc}(2) = 1$ ” survives nowhere on the right — and a prefix-free

insert has no way to reconstruct it.

The fix (carry the path). Give the insert the ancestor path of its anchor, $\text{Ins}(C, [1], 2)$. Now when 2 is dead, `resolve` climbs the carried path $[1]$ to the live 1, and both orders land 3 under 1 (Figure 5). The path is exactly the ancestry that the tombstone-free state threw away; carrying it on the operation restores commutation.




Figure 5: **The path restores convergence.** The carried path $[1]$ lets the late insert climb past the deleted 2 to the live 1, matching the other order. Mechanised as `prefix_converges`.

The slogan. *A tombstone keeps the deleted node’s parent in the **state**; the path keeps it on the **operation**.* Remove the tombstone and that parent has to ride along on any operation that might still need it — and on a single replica (no LCA) the insert is exactly such an operation. This is also why the two are *mutually exclusive but jointly unavoidable for the proof*: to satisfy the do-commutation VC you may drop the tombstone or drop the path, but not both. Dropping both defeats that route — provably (`RGA_PrefixFree_Impossible.lean`); the runnable witness that drops both yet still converges *via merge* is `RGA_PrefixFree_Impl.lean`.

4 The asymmetry: Ins needs the path, Del does not

Strikingly, only the *insert* needs the path. A delete can always recompute its target from the live state, because at the moment a delete runs, the node it reparents *to* (its target’s parent) is still present. Figure 6 contrasts the two: the delete reads its answer from a node that is alive *now*; the insert may need a node that was alive *then* but is gone *now*.

The intuition is temporal: a delete reparents *all current children* the instant it runs, keeping the forest connected; a concurrent (or reordered) insert is a child that was *not present* at that instant, so it missed the handoff and must supply its own ancestry.

5 What is mechanised

Every figure corresponds to a Lean fact in `Sal/MRDTs/RGA_Tombstone_Free/`:

- `prefix_converges` / `prefixfree_diverges` — Figures 4–5: with the path the pair commutes; drop it and the same accurate pair diverges. The *insert* path is necessary.
- `del_path_free` / `deldel_pathfree_converges` — Figure 6: the *delete* path is dispensable; a Del reads its reparent target $\text{anc}(\sigma, x)$ from the live state.




Figure 6: The asymmetry. Del hands off children to a parent that is still live at delete time, so it reads from state; a late Ins references a node that may have been spliced away, so it must carry the path. (*del_path_free* / *prefixfree_diverges*.)

- `rc_non_comm'` — the full commutation VC: on accurate, fresh-id, root-not-stored states, *every* operation pair commutes.
- `RGA_PrefixFree_Impossible.lean` — a parameterised theorem: no tombstone-free, prefix-free, local design can satisfy the `rc = Either` commutation VC.
- `RGA_PrefixFree_Impl.lean` — the runnable counterpart of Figures 3 and 4: its *merge* converges the concurrent case (LCA), yet its single-replica *do* diverges (*merge_converges_concurrent* vs. *single_replica_do_diverges*).

One honest caveat. The merge’s *climb* is bounded by a fuel equal to the node id, so its convergence relies on *id-monotone* anchors ($\text{anc}(t) < t$). That holds under monotone timestamp allocation but is *not* guaranteed by the forest shape alone; a non-monotone state can make *climb* run out of fuel under merge (*merge_breaks_wf* in `RGA_Reachability_Invariant.lean`). So Figure 3’s convergence is a theorem about reachable, monotone states — the precise extra condition the soundness argument must carry.